SE6014 · Shark-Tank Pitch

SHERLOCK

The threat hunter that turns a plain-English hunch into a **court-admissible verdict** — where every claim carries its receipt.

<1s	100%	0
hypothesis → verdict	claims evidence-backed	parsers written by hand

Group X · Members 1–5

THE PROBLEM

Threats move at machine speed. Hunting doesn’t.

In every SOC today

- Analysts **hand-load** logs and write a **one-off query** for every hunt.
- Hunts are slow, inconsistent, and walk out the door with the analyst.
- A senior hunter costs a fortune and can’t scale.

And when they try AI…

- LLMs **hallucinate** findings that were never in the data.
- You can’t tell a real detection from a confident guess.
- So nobody trusts it for a verdict that matters.

The trust gap is the whole problem. Speed is useless if you can’t believe the answer.

THE PRODUCT

You type a hypothesis. Sherlock returns a verdict — and proves it.

“An external actor brute-forced an admin account and established C2.”

↓ agents plan · query the SIEM through MCP · gather evidence ↓

🔴 **CONFIRMED — 95% confidence — 9 findings, every one cited.**

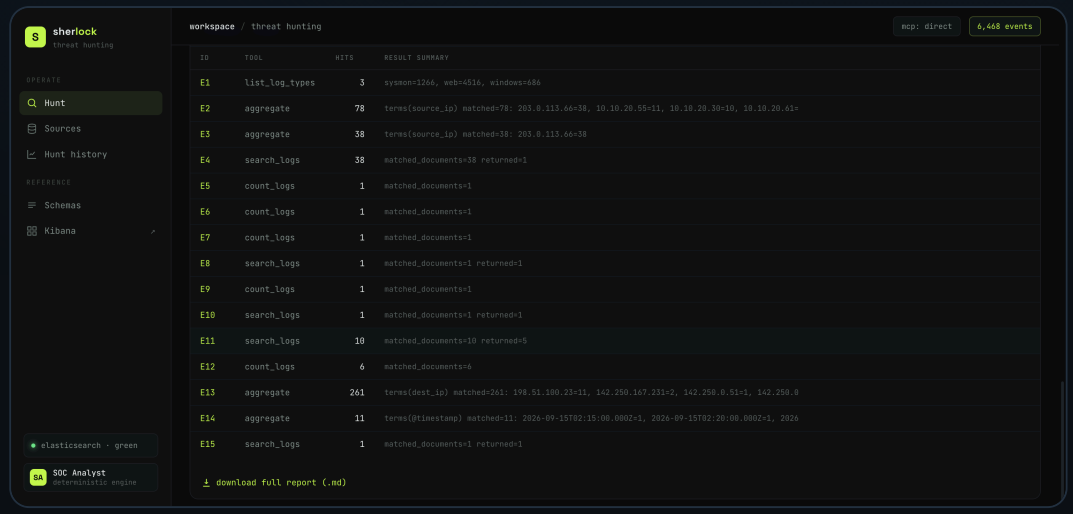
WHY YOU CAN TRUST IT · 1 OF 3

No evidence, no claim.

Every query the agent runs is written to an **evidence ledger** with the exact SIEM query and its hit count. A **verifier** then throws out any conclusion that isn't backed by a real result — **including a hallucinated citation** — *before* it reaches the report.

✓ Hallucination isn't discouraged. It's **mechanically blocked.**

✓ Every finding links to the query that proves it — reproducible by anyone.



WHY YOU CAN TRUST IT · 2 OF 3

Point it at any log source. Agents do the rest.

Zero-config ingestion

- An agent **inspects** the log, **matches** a parser template, or **builds one**, **validates** it, and **registers** it for reuse.
- Ingest by **pull**, **push**, or **stream** — same pipeline.
- We onboarded raw Windows, Sysmon, and plain-text web logs at a **100% parse rate**, with **no hand-written parsers.**

3
log types onboarded automatically

12
fields extracted from raw web text

WHY YOU CAN TRUST IT · 3 OF 3

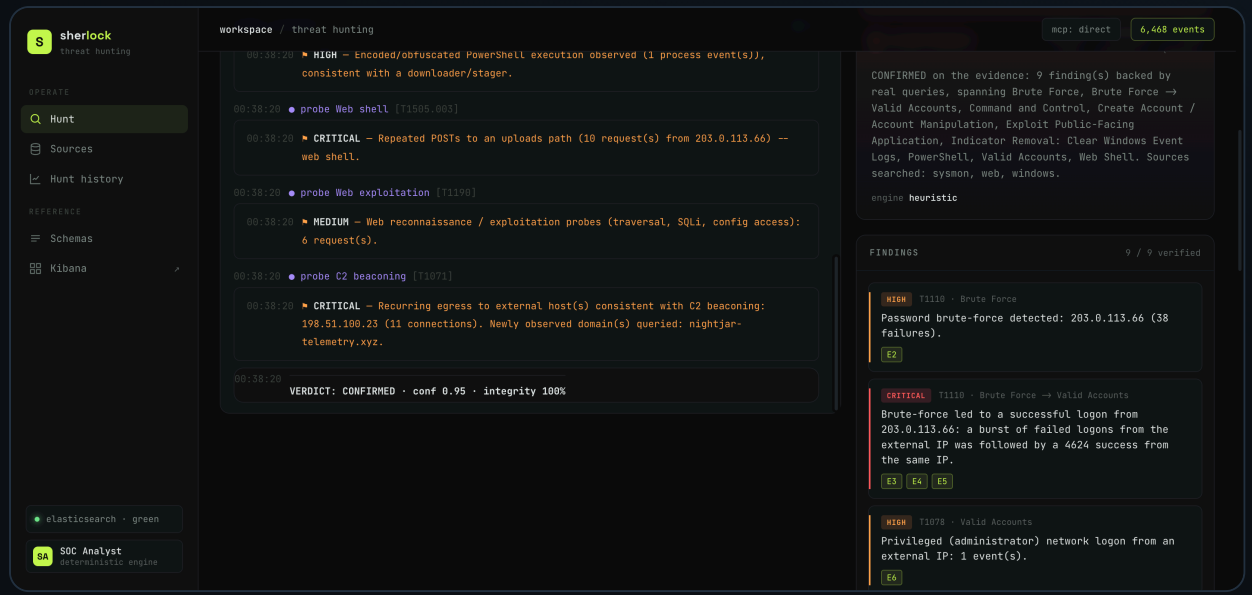
Same question, same answer — every time.

Run the identical hunt twice and the verdict, the findings, and the evidence **converge** — Jaccard **1.0**. A deterministic mode gives byte-identical, **auditable** runs for regulators and detection regression-testing. Repeatability is the precondition for trusting automation in a SOC — so we **measured it** and drove verdict variance to zero.

metric	run 1	run 2
verdict	confirmed	confirmed
findings	9	9
evidence integrity	100%	100%
finding-set overlap	**1.0**	

SEE IT WORK

Watch the agent hunt — live.



A full multi-stage intrusion — brute force → web shell → PowerShell → persistence → log-clearing → C2 — reconstructed across three sources in under a second.

✓ Every step streamed and cited.

UNDER THE HOOD

Agents that can only speak facts.



Built on the standard **ELK stack** and the **Model Context Protocol** — the same server plugs straight into Claude Desktop. Read-only tools bound the blast radius.

THE BUSINESS

A hunt capability every MSSP wants, none can staff.

Who buys

- Managed security service providers
- Enterprise SOC / threat-hunt teams
- Regulated & air-gapped shops (deterministic mode)

How we charge

- Per-analyst-seat SaaS
- Volume tier on data ingested
- On-prem licence for air-gapped & sovereign SOCs

The wedge

- Land as a **hunt copilot**
- Expand into **continuous** hunting
- Become the **detection-as-code** system of record

THE ASK

We built the trust layer for AI threat hunting. Help us take it to market.

Working	100%	2
end-to-end product, today	evidence integrity, proven	pilot MSSPs in the pipeline*

Seeking a design-partner MSSP and a seed round to harden for production and build continuous hunting.

*illustrative go-to-market target